

EDA and Prediction of US Accidents

Team: Jayatha Chandra, Avanti Dorle, Lavina Omprakash Talreja, Mansi Pravin

Thanki

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.metrics import precision_score, recall_score, f1_score
```

```
modeling_data = pd.read_csv("modeling_data.csv", na_values='')
display(modeling_data.head())
```

	Unnamed: 0	Severity	Start_Lat	Start_Lng	End_Lat	End_Lng	\
0	1	High	40.108910	-83.092860	40.112060	-83.031870	
1	2	Low	39.865420	-84.062800	39.865010	-84.048730	
2	3	Low	39.102660	-84.524680	39.102090	-84.523960	
3	4	Low	41.062130	-81.537840	41.062170	-81.535470	
4	5	High	39.172393	-84.492792	39.170476	-84.501798	

	Distance.mi.	Temperature.F.	Wind_Chill.F.	Humidity	...
Roundabout	\				
0	3.230	42.1	36.100000	58.0	...
1					
1	0.747	36.9	59.658231	91.0	...
1					
2	0.055	36.0	59.658231	97.0	...
1					
3	0.123	39.0	59.658231	55.0	...
1					
4	0.500	37.0	29.800000	93.0	...
1					

	Station	Stop	Traffic_Calming	Traffic_Signal	Turning_Loop	\
0	1	1	1	1	1	
1	1	1	1	1	1	
2	1	1	1	1	1	
3	1	1	1	1	1	
4	1	1	1	1	1	

	Sunrise_Sunset	Civil_Twilight	Nautical_Twilight
Astronomical_Twilight			
0	0	0	0
0			

1	0	0	0
0			
2	0	0	0
1			
3	0	0	1
1			
4	1	1	1
1			

[5 rows x 33 columns]

modeling_data.columns

```
Index(['Unnamed: 0', 'Severity', 'Start_Lat', 'Start_Lng', 'End_Lat',
      'End_Lng', 'Distance.mi.', 'Temperature.F.', 'Wind_Chill.F.',
      'Humidity', 'Pressure', 'Visibility', 'Wind_Direction',
      'Wind_Speed',
      'Precipitation', 'Weather_Condition', 'Amenity', 'Bump',
      'Crossing',
      'Give_Way', 'Junction', 'No_Exit', 'Railway', 'Roundabout',
      'Station',
      'Stop', 'Traffic_Calming', 'Traffic_Signal', 'Turning_Loop',
      'Sunrise_Sunset', 'Civil_Twilight', 'Nautical_Twilight',
      'Astronomical_Twilight'],
      dtype='object')
```

modeling_data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2845342 entries, 0 to 2845341
Data columns (total 33 columns):
#   Column                                Dtype
---  -
0   Unnamed: 0                            int64
1   Severity                              object
2   Start_Lat                             float64
3   Start_Lng                             float64
4   End_Lat                               float64
5   End_Lng                               float64
6   Distance.mi.                          float64
7   Temperature.F.                        float64
8   Wind_Chill.F.                         float64
9   Humidity                              float64
10  Pressure                               float64
11  Visibility                             float64
12  Wind_Direction                         int64
13  Wind_Speed                             float64
14  Precipitation                          float64
15  Weather_Condition                      int64
16  Amenity                                int64
17  Bump                                   int64
```

```
18 Crossing int64
19 Give_Way int64
20 Junction int64
21 No_Exit int64
22 Railway int64
23 Roundabout int64
24 Station int64
25 Stop int64
26 Traffic_Calming int64
27 Traffic_Signal int64
28 Turning_Loop int64
29 Sunrise_Sunset int64
30 Civil_Twilight int64
31 Nautical_Twilight int64
32 Astronomical_Twilight int64
dtypes: float64(12), int64(20), object(1)
memory usage: 716.4+ MB
```

```
modeling_data.dropna(axis=0)
modeling_data.isna().sum()
```

```
Unnamed: 0 0
Severity 0
Start_Lat 0
Start_Lng 0
End_Lat 0
End_Lng 0
Distance.mi. 0
Temperature.F. 0
Wind_Chill.F. 0
Humidity 0
Pressure 0
Visibility 0
Wind_Direction 0
Wind_Speed 0
Precipitation 0
Weather_Condition 0
Amenity 0
Bump 0
Crossing 0
Give_Way 0
Junction 0
No_Exit 0
Railway 0
Roundabout 0
Station 0
Stop 0
Traffic_Calming 0
Traffic_Signal 0
Turning_Loop 0
Sunrise_Sunset 0
```

```
Civil_Twilight          0
Nautical_Twilight       0
Astronomical_Twilight   0
dtype: int64
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
# Separate features and target variable
```

```
X = modeling_data.drop('Severity', axis=1)
```

```
y = modeling_data['Severity']
```

```
# Split the data into training, validation, and test sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
```

```
# Create a Random Forest classifier
```

```
rf_clf = RandomForestClassifier(n_estimators=100, random_state=42)
```

```
# Fit the classifier to the training data
```

```
rf_clf.fit(X_train, y_train)
```

```
# Make predictions on the testing data
```

```
y_pred_test = rf_clf.predict(X_test)
```

```
# Calculate accuracy of the classifier
```

```
accuracy = accuracy_score(y_test, y_pred_test)
```

```
print("Accuracy:", accuracy)
```

```
Accuracy: 0.9360218180923578
```

```
# Calculate accuracy, precision, recall, and F1-score on testing set
```

```
accuracy_test = accuracy_score(y_test, y_pred_test)
```

```
print("Accuracy:", accuracy_test)
```

```
# calculate confusion matrix
```

```
cm = confusion_matrix(y_test, y_pred_test)
```

```
print("Confusion Matrix:\n", cm)
```

```
# calculate classification report
```

```
cr = classification_report(y_test, y_pred_test)
```

```
print("Classification Report:\n", cr)
```

```
Accuracy: 0.9360218180923578
```

```
Confusion Matrix:
```

```
[[ 31232  25888]
```

```
 [ 10520 501429]]
```


```
Classification Report:
```

```
precision    recall  f1-score   support
```

High	0.75	0.55	0.63	57120
Low	0.95	0.98	0.96	511949
accuracy			0.94	569069
macro avg	0.85	0.76	0.80	569069
weighted avg	0.93	0.94	0.93	569069

```
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# create heatmap of confusion matrix
plt.figure(figsize=(8,6))
sns.heatmap(cm, annot=True, cmap='Blues')
plt.title('Confusion Matrix')
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.show()
```



```
accuracy_test = accuracy_score(y_test, y_pred_test)
print("Accuracy:", accuracy_test)
```

```

# calculate classification report
cr = classification_report(y_test, y_pred_test)
print("Classification Report:\n", cr)

# calculate macro-average precision, recall, and F1 score
macro_prec = precision_score(y_test, y_pred_test, average='macro')
macro_rec = recall_score(y_test, y_pred_test, average='macro')
macro_f1 = f1_score(y_test, y_pred_test, average='macro')
print("Macro-average Precision:", macro_prec)
print("Macro-average Recall:", macro_rec)
print("Macro-average F1 Score:", macro_f1)

```

Accuracy: 0.9360218180923578

Classification Report:

	precision	recall	f1-score	support
High	0.75	0.55	0.63	57120
Low	0.95	0.98	0.96	511949
accuracy			0.94	569069
macro avg	0.85	0.76	0.80	569069
weighted avg	0.93	0.94	0.93	569069

Macro-average Precision: 0.8494711066895253

Macro-average Recall: 0.7631148948096649

Macro-average F1 Score: 0.7983669535132445

```
rfc_scores = [accuracy_test, macro_prec, macro_rec, macro_f1]
```

```
# Import the necessary libraries
```

```

import pandas as pd
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from imblearn.under_sampling import RandomUnderSampler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score

```

```
# Read the data into a Pandas DataFrame
```

```
data = pd.read_csv('modeling_data.csv')
```

```
# Separate the features and target variable
```

```
X = data.drop(['Severity'], axis=1)
```

```
y = data['Severity']
```

```
# Undersample the majority class
```

```
rus = RandomUnderSampler(sampling_strategy=0.4)
```

```
X_resampled, y_resampled = rus.fit_resample(X, y)
```

```

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_resampled,
y_resampled, test_size=0.2, random_state=42)

# Scale the data using StandardScaler
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Set up the hyperparameter grid
hyperparameters = {
    'C': [0.001],
    'penalty': ['l2', 'none'],
    'solver': ['lbfgs', 'newton-cg']
}

# Perform grid search with cross-validation to find the best
hyperparameters
model = LogisticRegression(max_iter=3000)
grid_search = GridSearchCV(model, hyperparameters, cv=5, n_jobs=-1,
scoring='f1')
grid_search.fit(X_train_scaled, y_train)

# Make predictions on the testing data using the best model
y_pred = grid_search.predict(X_test_scaled)

# Evaluate the performance of the model using metrics such as
accuracy, precision, recall, and F1 score
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred, pos_label='High')
recall = recall_score(y_test, y_pred, pos_label='High')
f1 = f1_score(y_test, y_pred, pos_label='High')

print("Best hyperparameters:", grid_search.best_params_)
print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1 score:", f1)

/usr/local/lib/python3.9/dist-packages/sklearn/model_selection/
_search.py:952: UserWarning: One or more of the test scores are non-
finite: [nan nan nan nan]
  warnings.warn(

Best hyperparameters: {'C': 0.001, 'penalty': 'l2', 'solver': 'lbfgs'}
Accuracy: 0.7497068494927872
Precision: 0.6437644116152734
Recall: 0.2863799158524288
F1 score: 0.39641417483905905

```

```
# Calculate accuracy, precision, recall, and F1-score on testing set
accuracy_test = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy_test)
```

```
# calculate confusion matrix
cm = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:\n", cm)
```

```
# calculate classification report
cr = classification_report(y_test, y_pred)
print("Classification Report:\n", cr)
```

```
# calculate macro-average precision, recall, and F1 score
macro_prec = precision_score(y_test, y_pred, average='macro')
macro_rec = recall_score(y_test, y_pred, average='macro')
macro_f1 = f1_score(y_test, y_pred, average='macro')
print("Macro-average Precision:", macro_prec)
print("Macro-average Recall:", macro_rec)
print("Macro-average F1 Score:", macro_f1)
```

```
# calculate micro-average precision, recall, and F1 score
micro_prec = precision_score(y_test, y_pred, average='micro')
micro_rec = recall_score(y_test, y_pred, average='micro')
micro_f1 = f1_score(y_test, y_pred, average='micro')
print("Micro-average Precision:", micro_prec)
print("Micro-average Recall:", micro_rec)
print("Micro-average F1 Score:", micro_f1)
```

Accuracy: 0.7497068494927872

Confusion Matrix:

```
[[ 16472  41046]
 [  9115 133776]]
```

Classification Report:

	precision	recall	f1-score	support
High	0.64	0.29	0.40	57518
Low	0.77	0.94	0.84	142891
accuracy			0.75	200409
macro avg	0.70	0.61	0.62	200409
weighted avg	0.73	0.75	0.71	200409

Macro-average Precision: 0.7044885139381923

Macro-average Recall: 0.6112950170272075

Macro-average F1 Score: 0.619266345303217

Micro-average Precision: 0.7497068494927872

Micro-average Recall: 0.7497068494927872

Micro-average F1 Score: 0.7497068494927872

```
logistic_scores = [accuracy_test, macro_prec, macro_rec, macro_f1]
```

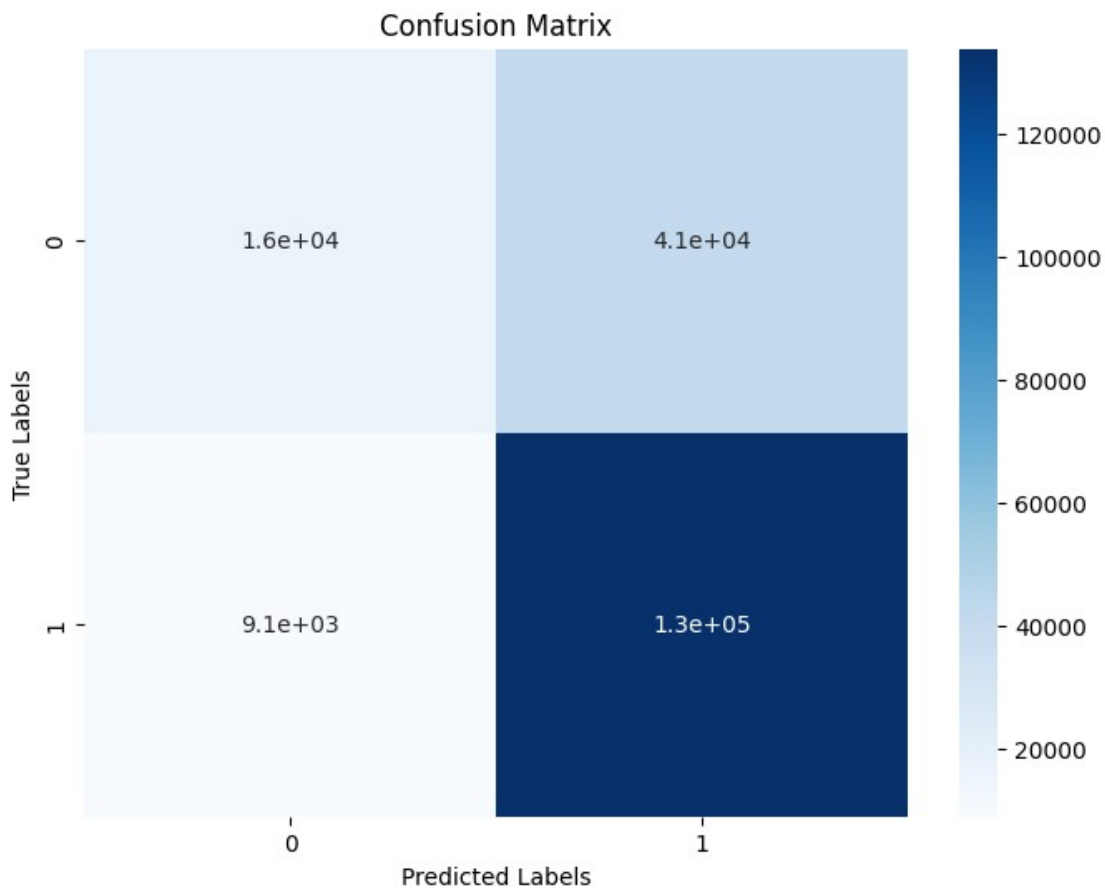


```
logistic_scores
```

```
[0.7497068494927872, 0.7044885139381923, 0.6112950170272075,  
0.619266345303217]
```

```
import matplotlib.pyplot as plt  
import seaborn as sns
```

```
# create heatmap of confusion matrix  
plt.figure(figsize=(8,6))  
sns.heatmap(cm, annot=True, cmap='Blues')  
plt.title('Confusion Matrix')  
plt.xlabel('Predicted Labels')  
plt.ylabel('True Labels')  
plt.show()
```



```
import numpy as np  
import matplotlib.pyplot as plt
```

```
labels = ['accuracy', 'macro_prec', 'macro_rec', 'macro_f1'] # replace  
with your own list of labels  
x = np.arange(len(labels)) # the label locations  
width = 0.35 # the width of the bars
```

```
fig, ax = plt.subplots()
rects1 = ax.bar(x - width/2, logistic_scores, width, label='Logistic
Regression')
rects2 = ax.bar(x + width/2, rfc_scores, width, label='Random Forest')

# Add some text for labels, title and custom x-axis tick labels, etc.
ax.set_ylabel('Score')
ax.set_title('Model Comparison')
ax.set_xticks(x)
ax.set_xticklabels(labels)
ax.legend()

fig.tight_layout()

plt.show()
```

